Contents

Preface **v**

Contents **xi**

1 The Deep Learning Revolution **1**
1.1 The Impact of Deep Learning 2
1.1.1 Medical diagnosis 2
1.1.2 Protein structure 3
1.1.3 Image synthesis 4
1.1.4 Large language models 5
1.2 A Tutorial Example 6
1.2.1 Synthetic data 6
1.2.2 Linear models 8
1.2.3 Error function 8
1.2.4 Model complexity 9
1.2.5 Regularization 12
1.2.6 Model selection 14
1.3 A Brief History of Machine Learning 16
1.3.1 Single-layer networks 17
1.3.2 Backpropagation 18
1.3.3 Deep networks 20

2 Probabilities **23**
2.1 The Rules of Probability 25
2.1.1 A medical screening example 25
2.1.2 The sum and product rules 26
2.1.3 Bayes' theorem 28
2.1.4 Medical screening revisited 30
2.1.5 Prior and posterior probabilities 31

2.1.6 Independent variables . . . 31
2.2 Probability Densities . . . 32
2.2.1 Example distributions . . . 33
2.2.2 Expectations and covariances . . . 34
2.3 The Gaussian Distribution . . . 36
2.3.1 Mean and variance . . . 37
2.3.2 Likelihood function . . . 37
2.3.3 Bias of maximum likelihood . . . 39
2.3.4 Linear regression . . . 40
2.4 Transformation of Densities . . . 42
2.4.1 Multivariate distributions . . . 44
2.5 Information Theory . . . 46
2.5.1 Entropy . . . 46
2.5.2 Physics perspective . . . 47
2.5.3 Differential entropy . . . 49
2.5.4 Maximum entropy . . . 50
2.5.5 Kullback–Leibler divergence . . . 51
2.5.6 Conditional entropy . . . 53
2.5.7 Mutual information . . . 54
2.6 Bayesian Probabilities . . . 54
2.6.1 Model parameters . . . 55
2.6.2 Regularization . . . 56
2.6.3 Bayesian machine learning . . . 57
Exercises . . . 58

3 Standard Distributions 65
3.1 Discrete Variables . . . 66
3.1.1 Bernoulli distribution . . . 66
3.1.2 Binomial distribution . . . 67
3.1.3 Multinomial distribution . . . 68
3.2 The Multivariate Gaussian . . . 70
3.2.1 Geometry of the Gaussian . . . 71
3.2.2 Moments . . . 74
3.2.3 Limitations . . . 75
3.2.4 Conditional distribution . . . 76
3.2.5 Marginal distribution . . . 79
3.2.6 Bayes' theorem . . . 81
3.2.7 Maximum likelihood . . . 84
3.2.8 Sequential estimation . . . 85
3.2.9 Mixtures of Gaussians . . . 86
3.3 Periodic Variables . . . 89
3.3.1 Von Mises distribution . . . 89
3.4 The Exponential Family . . . 94
3.4.1 Sufficient statistics . . . 97
3.5 Nonparametric Methods . . . 98

3.5.1 Histograms . . . 98
3.5.2 Kernel densities . . . 100
3.5.3 Nearest-neighbours . . . 103
Exercises . . . 105

4 Single-layer Networks: Regression 111
4.1 Linear Regression . . . 112
4.1.1 Basis functions . . . 112
4.1.2 Likelihood function . . . 114
4.1.3 Maximum likelihood . . . 115
4.1.4 Geometry of least squares . . . 117
4.1.5 Sequential learning . . . 117
4.1.6 Regularized least squares . . . 118
4.1.7 Multiple outputs . . . 119
4.2 Decision theory . . . 120
4.3 The Bias–Variance Trade-off . . . 123
Exercises . . . 128

5 Single-layer Networks: Classification 131
5.1 Discriminant Functions . . . 132
5.1.1 Two classes . . . 132
5.1.2 Multiple classes . . . 134
5.1.3 1-of-K coding . . . 135
5.1.4 Least squares for classification . . . 136
5.2 Decision Theory . . . 138
5.2.1 Misclassification rate . . . 139
5.2.2 Expected loss . . . 140
5.2.3 The reject option . . . 142
5.2.4 Inference and decision . . . 143
5.2.5 Classifier accuracy . . . 147
5.2.6 ROC curve . . . 148
5.3 Generative Classifiers . . . 150
5.3.1 Continuous inputs . . . 152
5.3.2 Maximum likelihood solution . . . 153
5.3.3 Discrete features . . . 156
5.3.4 Exponential family . . . 156
5.4 Discriminative Classifiers . . . 157
5.4.1 Activation functions . . . 158
5.4.2 Fixed basis functions . . . 158
5.4.3 Logistic regression . . . 159
5.4.4 Multi-class logistic regression . . . 161
5.4.5 Probit regression . . . 163
5.4.6 Canonical link functions . . . 164
Exercises . . . 166

6 Deep Neural Networks **171**
6.1 Limitations of Fixed Basis Functions . . . 172
6.1.1 The curse of dimensionality . . . 172
6.1.2 High-dimensional spaces . . . 175
6.1.3 Data manifolds . . . 176
6.1.4 Data-dependent basis functions . . . 178
6.2 Multilayer Networks . . . 180
6.2.1 Parameter matrices . . . 181
6.2.2 Universal approximation . . . 181
6.2.3 Hidden unit activation functions . . . 182
6.2.4 Weight-space symmetries . . . 185
6.3 Deep Networks . . . 186
6.3.1 Hierarchical representations . . . 187
6.3.2 Distributed representations . . . 187
6.3.3 Representation learning . . . 188
6.3.4 Transfer learning . . . 189
6.3.5 Contrastive learning . . . 191
6.3.6 General network architectures . . . 193
6.3.7 Tensors . . . 194
6.4 Error Functions . . . 194
6.4.1 Regression . . . 194
6.4.2 Binary classification . . . 196
6.4.3 multiclass classification . . . 197
6.5 Mixture Density Networks . . . 198
6.5.1 Robot kinematics example . . . 198
6.5.2 Conditional mixture distributions . . . 199
6.5.3 Gradient optimization . . . 201
6.5.4 Predictive distribution . . . 202
Exercises . . . 204

7 Gradient Descent **209**
7.1 Error Surfaces . . . 210
7.1.1 Local quadratic approximation . . . 211
7.2 Gradient Descent Optimization . . . 213
7.2.1 Use of gradient information . . . 214
7.2.2 Batch gradient descent . . . 214
7.2.3 Stochastic gradient descent . . . 214
7.2.4 Mini-batches . . . 216
7.2.5 Parameter initialization . . . 216
7.3 Convergence . . . 218
7.3.1 Momentum . . . 220
7.3.2 Learning rate schedule . . . 222
7.3.3 RMSProp and Adam . . . 223
7.4 Normalization . . . 224
7.4.1 Data normalization . . . 226

7.4.2 Batch normalization . . . 227
7.4.3 Layer normalization . . . 229
Exercises . . . 230

8 Backpropagation 233
8.1 Evaluation of Gradients . . . 234
8.1.1 Single-layer networks . . . 234
8.1.2 General feed-forward networks . . . 235
8.1.3 A simple example . . . 238
8.1.4 Numerical differentiation . . . 239
8.1.5 The Jacobian matrix . . . 240
8.1.6 The Hessian matrix . . . 242
8.2 Automatic Differentiation . . . 244
8.2.1 Forward-mode automatic differentiation . . . 246
8.2.2 Reverse-mode automatic differentiation . . . 249
Exercises . . . 250

9 Regularization 253
9.1 Inductive Bias . . . 254
9.1.1 Inverse problems . . . 254
9.1.2 No free lunch theorem . . . 255
9.1.3 Symmetry and invariance . . . 256
9.1.4 Equivariance . . . 259
9.2 Weight Decay . . . 260
9.2.1 Consistent regularizers . . . 262
9.2.2 Generalized weight decay . . . 264
9.3 Learning Curves . . . 266
9.3.1 Early stopping . . . 266
9.3.2 Double descent . . . 268
9.4 Parameter Sharing . . . 270
9.4.1 Soft weight sharing . . . 271
9.5 Residual Connections . . . 274
9.6 Model Averaging . . . 277
9.6.1 Dropout . . . 279
Exercises . . . 281

10 Convolutional Networks 287
10.1 Computer Vision . . . 288
10.1.1 Image data . . . 289
10.2 Convolutional Filters . . . 290
10.2.1 Feature detectors . . . 290
10.2.2 Translation equivariance . . . 291
10.2.3 Padding . . . 294
10.2.4 Strided convolutions . . . 294
10.2.5 Multi-dimensional convolutions . . . 295
10.2.6 Pooling . . . 296

10.2.7 Multilayer convolutions . . . 298
10.2.8 Example network architectures . . . 299
10.3 Visualizing Trained CNNs . . . 302
10.3.1 Visual cortex . . . 302
10.3.2 Visualizing trained filters . . . 303
10.3.3 Saliency maps . . . 305
10.3.4 Adversarial attacks . . . 306
10.3.5 Synthetic images . . . 308
10.4 Object Detection . . . 308
10.4.1 Bounding boxes . . . 309
10.4.2 Intersection-over-union . . . 310
10.4.3 Sliding windows . . . 311
10.4.4 Detection across scales . . . 313
10.4.5 Non-max suppression . . . 314
10.4.6 Fast region CNNs . . . 314
10.5 Image Segmentation . . . 315
10.5.1 Convolutional segmentation . . . 315
10.5.2 Up-sampling . . . 316
10.5.3 Fully convolutional networks . . . 318
10.5.4 The U-net architecture . . . 319
10.6 Style Transfer . . . 320
Exercises . . . 322

11 Structured Distributions 325
11.1 Graphical Models . . . 326
11.1.1 Directed graphs . . . 326
11.1.2 Factorization . . . 327
11.1.3 Discrete variables . . . 329
11.1.4 Gaussian variables . . . 332
11.1.5 Binary classifier . . . 334
11.1.6 Parameters and observations . . . 334
11.1.7 Bayes' theorem . . . 336
11.2 Conditional Independence . . . 337
11.2.1 Three example graphs . . . 338
11.2.2 Explaining away . . . 341
11.2.3 D-separation . . . 343
11.2.4 Naive Bayes . . . 344
11.2.5 Generative models . . . 346
11.2.6 Markov blanket . . . 347
11.2.7 Graphs as filters . . . 348
11.3 Sequence Models . . . 349
11.3.1 Hidden variables . . . 352
Exercises . . . 353

12 Transformers **357**
12.1 Attention . . . 358
12.1.1 Transformer processing . . . 360
12.1.2 Attention coefficients . . . 361
12.1.3 Self-attention . . . 362
12.1.4 Network parameters . . . 363
12.1.5 Scaled self-attention . . . 366
12.1.6 Multi-head attention . . . 366
12.1.7 Transformer layers . . . 368
12.1.8 Computational complexity . . . 370
12.1.9 Positional encoding . . . 371
12.2 Natural Language . . . 374
12.2.1 Word embedding . . . 375
12.2.2 Tokenization . . . 377
12.2.3 Bag of words . . . 378
12.2.4 Autoregressive models . . . 379
12.2.5 Recurrent neural networks . . . 380
12.2.6 Backpropagation through time . . . 381
12.3 Transformer Language Models . . . 382
12.3.1 Decoder transformers . . . 383
12.3.2 Sampling strategies . . . 386
12.3.3 Encoder transformers . . . 388
12.3.4 Sequence-to-sequence transformers . . . 390
12.3.5 Large language models . . . 390
12.4 Multimodal Transformers . . . 394
12.4.1 Vision transformers . . . 395
12.4.2 Generative image transformers . . . 396
12.4.3 Audio data . . . 399
12.4.4 Text-to-speech . . . 400
12.4.5 Vision and language transformers . . . 402
Exercises . . . 403

13 Graph Neural Networks **407**
13.1 Machine Learning on Graphs . . . 409
13.1.1 Graph properties . . . 410
13.1.2 Adjacency matrix . . . 410
13.1.3 Permutation equivariance . . . 411
13.2 Neural Message-Passing . . . 412
13.2.1 Convolutional filters . . . 413
13.2.2 Graph convolutional networks . . . 414
13.2.3 Aggregation operators . . . 416
13.2.4 Update operators . . . 418
13.2.5 Node classification . . . 419
13.2.6 Edge classification . . . 420
13.2.7 Graph classification . . . 420

13.3 General Graph Networks 420
13.3.1 Graph attention networks 421
13.3.2 Edge embeddings 421
13.3.3 Graph embeddings 422
13.3.4 Over-smoothing 422
13.3.5 Regularization 423
13.3.6 Geometric deep learning 424
Exercises 425

14 Sampling 429
14.1 Basic Sampling Algorithms 430
14.1.1 Expectations 430
14.1.2 Standard distributions 431
14.1.3 Rejection sampling 433
14.1.4 Adaptive rejection sampling 435
14.1.5 Importance sampling 437
14.1.6 Sampling-importance-resampling 439
14.2 Markov Chain Monte Carlo 440
14.2.1 The Metropolis algorithm 441
14.2.2 Markov chains 442
14.2.3 The Metropolis–Hastings algorithm 445
14.2.4 Gibbs sampling 446
14.2.5 Ancestral sampling 450
14.3 Langevin Sampling 451
14.3.1 Energy-based models 452
14.3.2 Maximizing the likelihood 453
14.3.3 Langevin dynamics 454
Exercises 456

15 Discrete Latent Variables 459
15.1 K-means Clustering 460
15.1.1 Image segmentation 464
15.2 Mixtures of Gaussians 466
15.2.1 Likelihood function 468
15.2.2 Maximum likelihood 470
15.3 Expectation–Maximization Algorithm 474
15.3.1 Gaussian mixtures 478
15.3.2 Relation to K-means 480
15.3.3 Mixtures of Bernoulli distributions 481
15.4 Evidence Lower Bound 485
15.4.1 EM revisited 486
15.4.2 Independent and identically distributed data 488
15.4.3 Parameter priors 489
15.4.4 Generalized EM 489
15.4.5 Sequential EM 490
Exercises 490

16 Continuous Latent Variables **495**
16.1 Principal Component Analysis . . . 497
16.1.1 Maximum variance formulation . . . 497
16.1.2 Minimum-error formulation . . . 499
16.1.3 Data compression . . . 501
16.1.4 Data whitening . . . 502
16.1.5 High-dimensional data . . . 504
16.2 Probabilistic Latent Variables . . . 506
16.2.1 Generative model . . . 506
16.2.2 Likelihood function . . . 507
16.2.3 Maximum likelihood . . . 509
16.2.4 Factor analysis . . . 513
16.2.5 Independent component analysis . . . 514
16.2.6 Kalman filters . . . 515
16.3 Evidence Lower Bound . . . 516
16.3.1 Expectation maximization . . . 518
16.3.2 EM for PCA . . . 519
16.3.3 EM for factor analysis . . . 520
16.4 Nonlinear Latent Variable Models . . . 522
16.4.1 Nonlinear manifolds . . . 522
16.4.2 Likelihood function . . . 524
16.4.3 Discrete data . . . 526
16.4.4 Four approaches to generative modelling . . . 527
Exercises . . . 527

17 Generative Adversarial Networks **533**
17.1 Adversarial Training . . . 534
17.1.1 Loss function . . . 535
17.1.2 GAN training in practice . . . 536
17.2 Image GANs . . . 539
17.2.1 CycleGAN . . . 539
Exercises . . . 544

18 Normalizing Flows **547**
18.1 Coupling Flows . . . 549
18.2 Autoregressive Flows . . . 552
18.3 Continuous Flows . . . 554
18.3.1 Neural differential equations . . . 554
18.3.2 Neural ODE backpropagation . . . 555
18.3.3 Neural ODE flows . . . 557
Exercises . . . 559

19 Autoencoders **563**
19.1 Deterministic Autoencoders . . . 564
19.1.1 Linear autoencoders . . . 564
19.1.2 Deep autoencoders . . . 565
19.1.3 Sparse autoencoders . . . 566
19.1.4 Denoising autoencoders . . . 567
19.1.5 Masked autoencoders . . . 567
19.2 Variational Autoencoders . . . 569
19.2.1 Amortized inference . . . 572
19.2.2 The reparameterization trick . . . 574
Exercises . . . 578

20 Diffusion Models **581**
20.1 Forward Encoder . . . 582
20.1.1 Diffusion kernel . . . 583
20.1.2 Conditional distribution . . . 584
20.2 Reverse Decoder . . . 585
20.2.1 Training the decoder . . . 587
20.2.2 Evidence lower bound . . . 588
20.2.3 Rewriting the ELBO . . . 589
20.2.4 Predicting the noise . . . 591
20.2.5 Generating new samples . . . 592
20.3 Score Matching . . . 594
20.3.1 Score loss function . . . 595
20.3.2 Modified score loss . . . 596
20.3.3 Noise variance . . . 597
20.3.4 Stochastic differential equations . . . 598
20.4 Guided Diffusion . . . 599
20.4.1 Classifier guidance . . . 600
20.4.2 Classifier-free guidance . . . 600
Exercises . . . 603

Appendix A Linear Algebra **609**
A.1 Matrix Identities . . . 609
A.2 Traces and Determinants . . . 610
A.3 Matrix Derivatives . . . 611
A.4 Eigenvectors . . . 612

Appendix B Calculus of Variations **617**

Appendix C Lagrange Multipliers **621**

Bibliography **625**

Index **641**